

Runbook RAG – Dev & Prod Environment Guide

This document defines the standard process for managing Runbook PDF documents used in Retrieval-Augmented Generation (RAG) across Development and Production environments. This file is intended to be committed under the DOCS folder and treated as a living reference.

1. Purpose

To ensure consistent, auditable, and safe ingestion of runbook PDFs into the RAG pipeline without impacting MCP workflows, UI components, or production stability.

2. Folder Structure (Mandatory)

All runbook documents must reside under the DOCS directory. This structure must be identical for Dev and Prod environments.

- DOCS/runbooks/azure/
- DOCS/runbooks/aws/
- DOCS/runbooks/general/

3. Environment Separation

Dev and Prod environments must remain logically isolated. Documents may be shared, but vector stores, embeddings, and ingestion jobs must be environment-specific.

- DEV: Used for testing new or updated runbooks before promotion
- PROD: Serves live queries from approved and validated runbooks only

4. Adding a New Runbook (Dev)

- Place PDF under DOCS/runbooks//
- Validate PDF readability and naming
- Ingest into Dev vector store
- Test with sample RAG queries

5. Promotion to Production

- Confirm metadata completeness
- Approve version and content
- Re-ingest PDF into Prod vector store
- Do not copy vectors between environments

6. Versioning Rules

- Minor updates increment version number
- Major changes require new document ID
- Old versions may be retained but marked deprecated

7. What Does NOT Change

- UI components (Home, MCP, Agents pages)
- MCP workflow definitions
- API contracts

8. Compliance & Audit Notes

DOCS is treated as the authoritative knowledge boundary. All changes are traceable via Git history and environment-specific ingestion logs.

End of Document – Commit this file under DOCS/ for reference.