

Runbook RAG – Ingestion Architecture

Purpose

Defines the ingestion architecture for converting PDF runbooks into searchable RAG knowledge.

High-Level Flow

- DOCS/runbooks → PDF Loader
- Text Normalization
- Chunking + Metadata Attachment
- Embeddings Generation
- Vector Store Persistence

Design Principles

- Idempotent ingestion
- Environment isolation
- Append-only vectors
- Metadata-first design

End of Document – Commit under DOCS/