

LLM-SRE Site: Directory Structure, Wiring Diagram, and Threat Model

This document captures a practical reference architecture for the llm-sre-site project: how the repo is organized, how UI/infra/backend are wired in AWS, and a lightweight threat model focused on the /api/runbooks/ask RAG flow.

1) Recommended Repository Directory Structure

```

llm-sre-site/
  ui/                                     # Vite + React frontend (static)
    src/
      components/
        AskRunbooks.jsx                  # RAG UI (question -> /api/runbooks/ask -> answer + sou
        pages/
          Agents.jsx
          Rag.jsx
          Runbooks.jsx
        lib/
          api.js
          App.jsx
        public/
        package.json
        vite.config.js

  services/
    agent_api/
      app.py
      requirements.txt
      Dockerfile
      # Lambda handler (agents + runbooks/ask)
      # openai + chroma + sqlite shim deps
      # Lambda container image build

  modules/
    agent_api/
      main.tf
      variables.tf
      outputs.tf
      # Lambda (Image), API Gateway v2 routes, IAM, CORS

    scripts/
      index_runbooks_chroma.py          # builds local chroma_store/ from PDFs + embeddings
      (optional) runbook_manifest.py   # incremental diffing + dry-run (hashetag)
      ...

    infra/                                # root Terraform (calls modules/*)
      main.tf
      env/dev.tfvars
      ...
      ...

  README.md

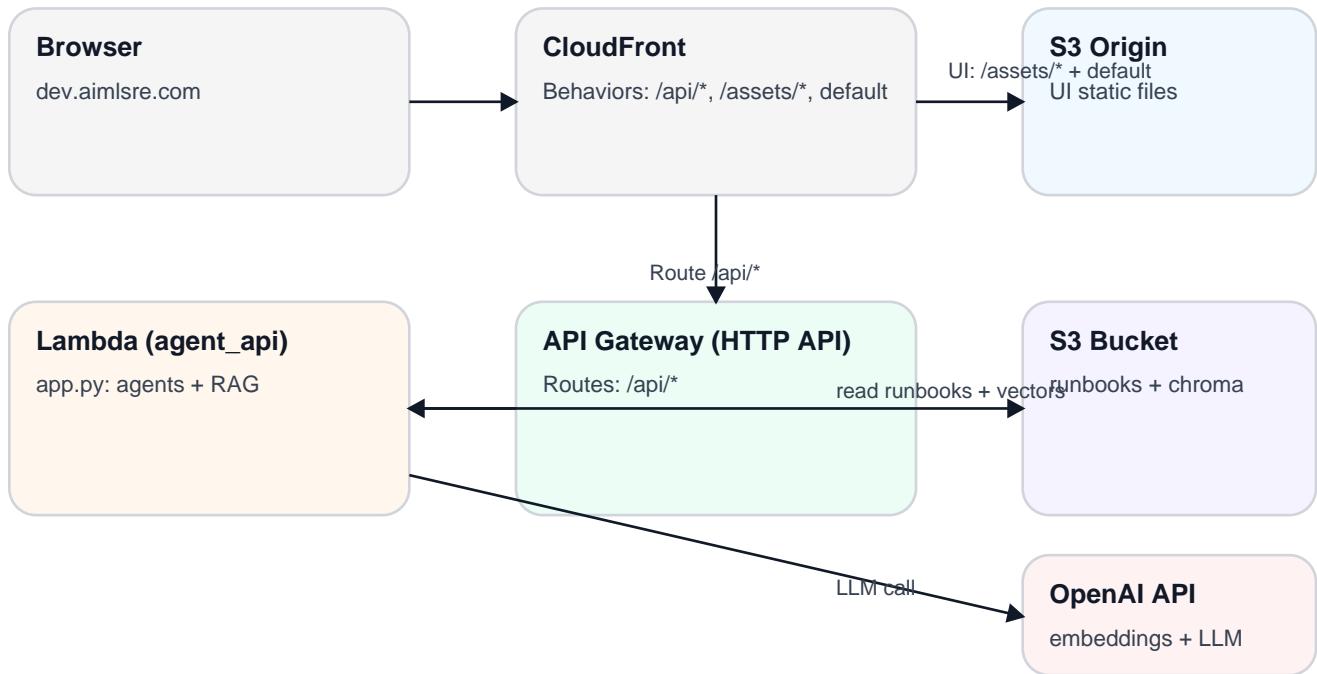
```

Notes

- Keep secrets (OpenAI key) out of the repo. Use Lambda environment variables or AWS Secrets Manager.
- Keep the Chroma index as an artifact in S3 under **knowledge/vectors/<env>/chroma/** and download to **/tmp** at runtime.
- scripts/ is a good location for the indexing job; run it locally or in CI, then upload the resulting chroma_store/ to S3.

2) End-to-End Wiring: UI -> CloudFront -> API Gateway -> Lambda -> S3/OpenAI

UI is served from an S3 static site origin behind CloudFront. API traffic is routed to an API Gateway v2 HTTP API origin. Lambda handles /api/agents, /api/agent/run, and /api/runbooks/ask (RAG).



Key AWS resources

Layer	Primary resource	Purpose / Notes
Edge	CloudFront Distribution	Routes UI + API; ordered behaviors: <code>/assets/*</code> -> S3, <code>/api/*</code> -> API GW, <code>default</code> -> S3
UI origin	S3 bucket (static website)	Hosts Vite build output (dist/)
API	API Gateway v2 (HTTP API)	Routes: GET <code>/api/health</code> , GET <code>/api/agents</code> , POST <code>/api/agent/run</code> , POST <code>/api/runbook/ask</code>
Compute	Lambda (container image)	services/agent_api Docker image includes chromadb + sqlite shim
Knowledge	S3 bucket (runbooks + vectors)	knowledge/runbooks/*.pdf + knowledge/vectors/<env>/chroma/*
LLM	OpenAI API	Embeddings + Responses API for answer generation

3) RAG Flow (Runbooks Q&A;)

1. User enters a question in the UI (AskRunbooks component).
2. UI POSTs JSON to /api/runbooks/ask via CloudFront -> API Gateway -> Lambda.
3. Lambda downloads Chroma store from S3 to /tmp/chroma_store (cold start or cache-miss).
4. Lambda creates an embedding for the question (OpenAI embeddings).
5. Lambda queries Chroma for top_k chunks and generates an answer using only the excerpts.
6. UI renders answer + sources (file + chunk).

Mermaid sequence diagram (copy into your docs/wiki)

```
```mermaid
sequenceDiagram
 autonumber
 participant U as User (Browser)
 participant CF as CloudFront
 participant APIGW as API Gateway (HTTP API)
 participant L as Lambda (agent_api)
 participant S3 as S3 (runbooks + vectors)
 participant OAI as OpenAI API

 U->>CF: POST /api/runbooks/ask {question, top_k}
 CF->>APIGW: Forward /api/* behavior
 APIGW-->L: Invoke Lambda proxy
 L-->>S3: Download vectors prefix to /tmp (if needed)
 L-->>OAI: Create embedding(question)
 L-->>L: Chroma query top_k
 L-->>OAI: Generate answer from excerpts
 L-->>APIGW: {answer, sources[]}
 APIGW-->>CF: Response
 CF-->>U: Render answer in UI
```
```

4) Threat Model (Practical, lightweight)

Scope: public site + API endpoints. Focus on protecting credentials, controlling cost, and preventing data exfiltration or prompt injection via untrusted content.

| Threat | Example | Mitigation (recommended) |
|-------------------------------|--|---|
| Secrets leakage | OpenAI API key exposed in repo/UI, logs, or client-side code. | Keep OPENAI_API_KEY server-side only. Store in Lambda env or Secrets Manager; never bundle into UI. Scrub logs. |
| Prompt injection via runbooks | Malicious content inside PDFs attempts to override instructions. | Use strict system prompt: answer only from excerpts. Strip/limit context length; consider corpus allowlist. |
| Data exfiltration | Model outputs sensitive runbook details to unauthorized users. | Add AuthN/AuthZ if needed (Cognito/JWT). Rate limit/WAF. Consider redaction for secrets in runbooks. |
| Cost abuse / DoS | Attackers spam /api/runbooks/ask to burn tokens and Lambda time. | Throttle (API GW), WAF rate-based rules, top_k cap, max tokens, and per-IP limits. |
| S3 abuse / path traversal | Attacker forces Lambda to download unexpected S3 objects. | Do not accept arbitrary S3 keys from client for RAG. Keep vectors prefix fixed in env; validate inputs. |
| Supply chain risks | Compromised dependencies in container image. | Pin versions, scan ECR images, enable Dependabot, and restrict outbound egress where possible. |

Mermaid data-flow diagram (DFD-lite)

```
```mermaid
flowchart LR
 U[User Browser] -->|HTTPS| CF[CloudFront]
 CF -->|/assets/*, default| S3UI[S3 UI Bucket]
 CF -->|/api/*| APIGW[API Gateway HTTP API]
 APIGW --> L[Lambda agent_api]
 L -->|GetObject/List| S3KB[S3 Knowledge Bucket]
 L -->|Embeddings + Responses| OAI[OpenAI API]
```

```

Next hardening steps: add authentication (Cognito), restrict origins, enable WAF rate limiting, add request size limits, and optionally move OpenAI key into Secrets Manager with rotation.